

Assignment 4: Data Wrangling (Fall 2024)

Tori Newton

OVERVIEW

This exercise accompanies the lessons in Environmental Data Analytics on Data Wrangling

Directions

1. Rename this file `<FirstLast>_A04_DataWrangling.Rmd` (replacing `<FirstLast>` with your first and last name).
2. Change “Student Name” on line 3 (above) with your name.
3. Work through the steps, **creating code and output** that fulfill each instruction.
4. Be sure to **answer the questions** in this assignment document.
5. When you have completed the assignment, **Knit** the text and code into a single PDF file.
6. Ensure that code in code chunks does not extend off the page in the PDF.

Set up your session

- 1a. Load the `tidyverse`, `lubridate`, and `here` packages into your session.
 - 1b. Check your working directory.
 - 1c. Read in all four raw data files associated with the EPA Air dataset, being sure to set string columns to be read in a factors. See the README file for the EPA air datasets for more information (especially if you have not worked with air quality data previously).
2. Add the appropriate code to reveal the dimensions of the four datasets.

```
#1a: loading packages
```

```
library(tidyverse)
```

```
library(lubridate)
```

```
library(here)
```

```
#1b: check working directory
```

```
here()
```

```
## [1] "/home/guest/EDA_Spring2025"
```

```
#1c: read in raw data files
```

```
#1: EPAair_03_NC2018
```

```
EPAair_03_NC2018 <- read.csv(  
  file=here("Data/Raw/EPAair_03_NC2018_raw.csv"),  
  stringsAsFactors = TRUE
```

```

)

#2: EPAair_03_NC2019
EPAair_03_NC2019 <- read.csv(
  file=here("Data/Raw/EPAair_03_NC2019_raw.csv"),
  stringsAsFactors = TRUE
)

#3: EPAair_PM25_NC2018
EPAair_PM25_NC2018 <- read.csv(
  file=here("Data/Raw/EPAair_PM25_NC2018_raw.csv"),
  stringsAsFactors = TRUE
)

#4: EPAair_PM25_NC2019
EPAair_PM25_NC2019 <- read.csv(
  file=here("Data/Raw/EPAair_PM25_NC2019_raw.csv"),
  stringsAsFactors = TRUE
)

#2: dimensions of four datasets
dim(EPAair_03_NC2018)

```

```
## [1] 9737    20
```

```
dim(EPAair_03_NC2019)
```

```
## [1] 10592    20
```

```
dim(EPAair_PM25_NC2018)
```

```
## [1] 8983    20
```

```
dim(EPAair_PM25_NC2019)
```

```
## [1] 8581    20
```

All four datasets should have the same number of columns but unique record counts (rows). Do your datasets follow this pattern?

Answer: Yes my datasets follow this pattern. They all have 20 columns but unique rows.

Wrangle individual datasets to create processed files.

3. Change the Date columns to be date objects.
4. Select the following columns: Date, DAILY_AQI_VALUE, Site.Name, AQS_PARAMETER_DESC, COUNTY, SITE_LATITUDE, SITE_LONGITUDE
5. For the PM2.5 datasets, fill all cells in AQS_PARAMETER_DESC with “PM2.5” (all cells in this column should be identical).

6. Save all four processed datasets in the Processed folder. Use the same file names as the raw files but replace “raw” with “processed”.

```
#3: change Date columns to be date objects
EPAair_03_NC2018$Date <- mdy(EPAair_03_NC2018$Date)
EPAair_03_NC2019$Date <- mdy(EPAair_03_NC2019$Date)
EPAair_PM25_NC2018$Date <- mdy(EPAair_PM25_NC2018$Date)
EPAair_PM25_NC2019$Date <- mdy(EPAair_PM25_NC2019$Date)

#4: select the requested columns
EPAair_03_NC2018_select <- EPAair_03_NC2018 %>%
  select(Date, DAILY_AQI_VALUE, Site.Name, AQS_PARAMETER_DESC, COUNTY,
    SITE_LATITUDE, SITE_LONGITUDE)
EPAair_03_NC2019_select <- EPAair_03_NC2019 %>%
  select(Date, DAILY_AQI_VALUE, Site.Name, AQS_PARAMETER_DESC, COUNTY,
    SITE_LATITUDE, SITE_LONGITUDE)
EPAair_PM25_NC2018_select <- EPAair_PM25_NC2018 %>%
  select(Date, DAILY_AQI_VALUE, Site.Name, AQS_PARAMETER_DESC, COUNTY,
    SITE_LATITUDE, SITE_LONGITUDE)
EPAair_PM25_NC2019_select <- EPAair_PM25_NC2019 %>%
  select(Date, DAILY_AQI_VALUE, Site.Name, AQS_PARAMETER_DESC, COUNTY,
    SITE_LATITUDE, SITE_LONGITUDE)

#5: fill all cells in AQS_PARAMETER_DESC with "PM2.5"
EPAair_PM25_NC2018_PM2.5 <- EPAair_PM25_NC2018_select %>%
  mutate(AQS_PARAMETER_DESC = "PM2.5")
EPAair_PM25_NC2019_PM2.5 <- EPAair_PM25_NC2019_select %>%
  mutate(AQS_PARAMETER_DESC = "PM2.5")

#6: save all four processed datasets in the Processed folder
write.csv(EPAair_03_NC2018_select, row.names = FALSE,
  file = "./Data/Processed/EPAair_03_NC2018_Processed.csv")
write.csv(EPAair_03_NC2019_select, row.names = FALSE,
  file = "./Data/Processed/EPAair_03_NC2019_Processed.csv")
write.csv(EPAair_PM25_NC2018_PM2.5, row.names = FALSE,
  file = "./Data/Processed/EPAair_PM25_NC2018_Processed.csv")
write.csv(EPAair_PM25_NC2019_PM2.5, row.names = FALSE,
  file = "./Data/Processed/EPAair_PM25_NC2019_Processed.csv")
```

Combine datasets

7. Combine the four datasets with `rbind`. Make sure your column names are identical prior to running this code.
8. Wrangle your new dataset with a pipe function (`%>%`) so that it fills the following conditions:
 - Include only sites that the four data frames have in common:

“Linville Falls”, “Durham Armory”, “Leggett”, “Hattie Avenue”,
“Clemmons Middle”, “Mendenhall School”, “Frying Pan Mountain”, “West Johnston Co.”, “Garinger High School”, “Castle Hayne”, “Pitt Agri. Center”, “Bryson City”, “Millbrook School”

(the function `intersect` can figure out common factor levels - but it will include sites with missing site information, which you don’t want...)

- Some sites have multiple measurements per day. Use the split-apply-combine strategy to generate daily means: group by date, site name, AQS parameter, and county. Take the mean of the AQI value, latitude, and longitude.
 - Add columns for “Month” and “Year” by parsing your “Date” column (hint: lubridate package)
 - Hint: the dimensions of this dataset should be 14,752 x 9.
9. Spread your datasets such that AQI values for ozone and PM2.5 are in separate columns. Each location on a specific date should now occupy only one row.
 10. Call up the dimensions of your new tidy dataset.
 11. Save your processed dataset with the following file name: “EPAair_O3_PM25_NC1819_Processed.csv”

```
#7: combine the four datasets
Datasets_Processed_Combined <- rbind(EPAair_O3_NC2018_select,
                                     EPAair_O3_NC2019_select,
                                     EPAair_PM25_NC2018_PM2.5,
                                     EPAair_PM25_NC2019_PM2.5)

#8: wrangle the dataset with a pipe function
#define common sites
common_sites <- c("Linville Falls", "Durham Armory", "Leggett", "Hattie Avenue",
                  "Clemmons Middle", "Mendenhall School", "Frying Pan Mountain",
                  "West Johnston Co.", "Garinger High School", "Castle Hayne",
                  "Pitt Agri. Center", "Bryson City", "Millbrook School")

New_Tidy_dataset <- Datasets_Processed_Combined %>%
  filter(Site.Name %in% common_sites) %>%
  group_by(Date, Site.Name, AQS_PARAMETER_DESC, COUNTY) %>%
  summarize(
    Mean_AQI = mean(DAILY_AQI_VALUE),
    Mean_Latitude = mean(SITE_LATITUDE),
    Mean_Longitude = mean(SITE_LONGITUDE)
  ) %>%
  mutate(
    Month = month(Date),
    Year = year(Date)
  )
```

```
## 'summarise()' has grouped output by 'Date', 'Site.Name', 'AQS_PARAMETER_DESC'.
## You can override using the '.groups' argument.
```

```
#check dimensions of new dataset
dim(New_Tidy_dataset) #The dimensions are 14752 x 9.
```

```
## [1] 14752      9
```

```
#9: spread datasets so AQI values for ozone and PM2.5 are in separate columns
spread_datasets <- pivot_wider(New_Tidy_dataset,
                               names_from = AQS_PARAMETER_DESC,
                               values_from = Mean_AQI)

#10: dimensions of new tidy dataset
dim(spread_datasets)
```

```
## [1] 8976    9
```

```
#11: save processed dataset
write.csv(spread_datasets, row.names = FALSE,
          file = "./Data/Processed/EPAair_03_PM25_NC1819_Processed.csv")
```

Generate summary tables

12. Use the split-apply-combine strategy to generate a summary data frame. Data should be grouped by site, month, and year. Generate the mean AQI values for ozone and PM2.5 for each group. Then, add a pipe to remove instances where mean **ozone** values are not available (use the function **drop_na** in your pipe). It's ok to have missing mean PM2.5 values in this result.

13. Call up the dimensions of the summary dataset.

```
#12: generate a summary data frame
summary_df <- spread_datasets %>%
  group_by(Site.Name, Month, Year) %>%
  summarize(
    Mean_AQI_PM2.5 = mean(PM2.5),
    Mean_AQI_Ozone = mean(Ozone),
  ) %>%
  drop_na (Mean_AQI_Ozone)
```

```
## 'summarise()' has grouped output by 'Site.Name', 'Month'. You can override
## using the '.groups' argument.
```

```
#13: dimensions for summary dataset
dim(summary_df)
```

```
## [1] 182    5
```

14. Why did we use the function **drop_na** rather than **na.omit**? Hint: replace **drop_na** with **na.omit** in part 12 and observe what happens with the dimensions of the summary data frame.

Answer: The function 'drop_na' allows us to selectively remove rows where only mean ozone is NA, while keeping rows where mean PM2.5 is NA. 'na.omit' removes any row that has at least one NA in any column. This means that if mean PM2.5 is missing but mean ozone is available, the entire row would be dropped which would result in a smaller summary data frame than intended.